

Importing libraries and datasets

In [1]:

```
import numpy as np
import pandas as pd
import datetime as dt

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
import xgboost as xgb

from sklearn.metrics import precision_score, recall_score, f1_score, auc, roc_auc_score, a
from xgboost import plot_importance
```

In [2]:

```
df_response = pd.read_csv('Retail_Data_Response.csv')
df_transactions = pd.read_csv('Retail_Data_Transactions.csv', parse_dates=['trans_date'])
```

In [3]:

```
df_response.head()
```

Out[3]:

	customer_id	response
0	CS1112	0
1	CS1113	0
2	CS1114	1
3	CS1115	1
4	CS1116	1

In [4]:

```
df_transactions.head()
```

Out[4]:

	customer_id	trans_date	tran_amount
0	CS5295	2013-02-11	35
1	CS4768	2015-03-15	39
2	CS2122	2013-02-26	52
3	CS1217	2011-11-16	99
4	CS1850	2013-11-20	78

In [5]:

```
print(df_transactions['trans_date'].min())
print(df_transactions['trans_date'].max())
```

```
2011-05-16 00:00:00
2015-03-16 00:00:00
```

Data Preparation

```
In [6]: ## since the last date of the data is 16 March 2015, the campaign date is assumed to be 17 March 2015

campaign_date = dt.datetime(2015,3,17)
df_transactions['recent'] = campaign_date - df_transactions['trans_date']
df_transactions['recent'].astype('timedelta64[D]')
df_transactions['recent'] = df_transactions['recent'] / np.timedelta64(1, 'D')
df_transactions.head()
```

```
Out[6]:
```

	customer_id	trans_date	tran_amount	recent
0	CS5295	2013-02-11	35	764.0
1	CS4768	2015-03-15	39	2.0
2	CS2122	2013-02-26	52	749.0
3	CS1217	2011-11-16	99	1217.0
4	CS1850	2013-11-20	78	482.0

```
In [7]: ## create data set with CLV variables

df_clv = df_transactions.groupby('customer_id').agg({'recent': lambda x: x.min(),
                                                    'customer_id': lambda x: len(x),
                                                    'tran_amount': lambda x: x.sum(),
                                                    'trans_date': lambda x: (x.max() - x.min()).days})

df_clv.rename(columns={'recent': 'recency',
                      'customer_id': 'frequency',
                      'tran_amount': 'monetary_value',
                      'trans_date': 'AOU'}, inplace=True)

df_clv['avg_size'] = df_clv['monetary_value'] / df_clv['frequency']
```

```
In [8]: df_clv = df_clv.reset_index()
df_clv.head()
```

```
Out[8]:
```

	customer_id	recency	frequency	monetary_value	AOU	avg_size
0	CS1112	62.0	15	1012	1309	67.466667
1	CS1113	36.0	20	1490	1354	74.500000
2	CS1114	33.0	19	1432	1309	75.368421
3	CS1115	12.0	22	1659	1303	75.409091
4	CS1116	204.0	13	857	1155	65.923077

Calculating response rate

```
In [9]: response_rate = df_response.groupby('response').agg({'customer_id': lambda x: len(x)}).reset_index()
response_rate.head()
```

```
Out[9]:
```


	response	customer_id
0	0	6237

	response	customer_id
1	1	647

In [10]:

```
plt.figure(figsize=(5,5))
x=range(2)
plt.bar(x,response_rate['customer_id'])
plt.xticks(response_rate.index)
plt.title('Response Distribution')
plt.xlabel('Convert or Not')
plt.ylabel('no. of users')
plt.show()

## data is imbalanced
```



In [11]:

```
## merging two data sets - CLV

df_modeling_clv = pd.merge(df_response,df_clv)
df_modeling_clv.head()
```

Out[11]:

	customer_id	response	recency	frequency	monetary_value	AOU	avg_size
0	CS1112	0	62.0	15	1012	1309	67.466667
1	CS1113	0	36.0	20	1490	1354	74.500000
2	CS1114	1	33.0	19	1432	1309	75.368421
3	CS1115	1	12.0	22	1659	1303	75.409091
4	CS1116	1	204.0	13	857	1155	65.923077

Creating train and test dataset


In [12]:

```
## splitting dataframe into X and y


X_clv = df_modeling_clv.drop(columns=['response','customer_id'])
y_clv = df_modeling_clv['response']
```

```
In [13]: ## creating train and test dataset
X_train_clv, X_test_clv, y_train_clv, y_test_clv = train_test_split(X_clv, y_clv, test_size=0.2)
```


```
In [14]: for i, col_i in enumerate(df_modeling_clv[['recency', 'frequency', 'monetary_value', 'AOU']]):
    for j, col_j in enumerate(df_modeling_clv[['recency', 'frequency', 'monetary_value', 'AOU']]):
        if i < j:
            plt.title(col_i + ' and ' + col_j)
            sns.scatterplot(data=df_modeling_clv, x=col_i, y=col_j, hue='response')
            sns.despine()
            plt.show()
```




recency and AOU




recency and avg_size




frequency and monetary_value




frequency and AOU

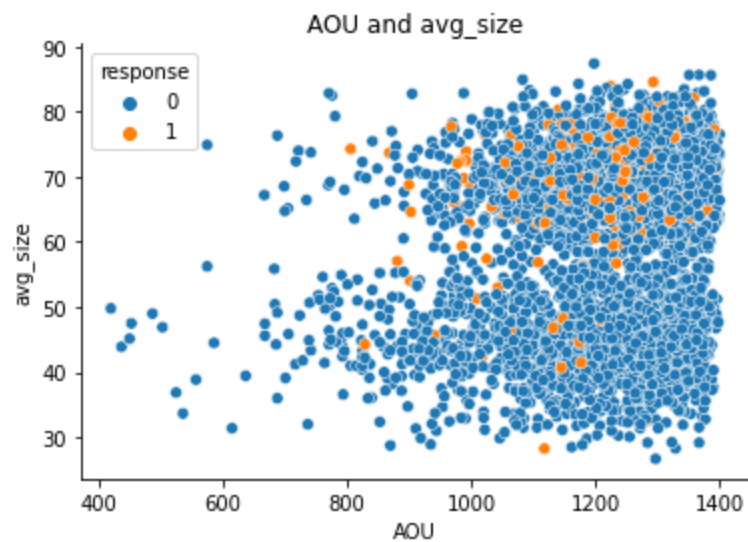



frequency and avg_size



monetary_value and AOU





```
In [15]: corrM = df_modeling_clv.corr()
corrM
```

```
Out[15]:
```

	response	recency	frequency	monetary_value	AOU	avg_size
response	1.000000	0.023405	0.201333	0.205543	0.002463	0.142133
recency	0.023405	1.000000	-0.286019	-0.270120	-0.715406	-0.176771
frequency	0.201333	-0.286019	1.000000	0.934468	0.401191	0.535893
monetary_value	0.205543	-0.270120	0.934468	1.000000	0.376935	0.786731
AOU	0.002463	-0.715406	0.401191	0.376935	1.000000	0.245546
avg_size	0.142133	-0.176771	0.535893	0.786731	0.245546	1.000000

Fixing imbalanced with SMOTE

```
In [16]: sm = SMOTE(random_state=0)

sm.fit(X_train_clv, y_train_clv)
X_SMOTE_clv, y_SMOTE_clv = sm.fit_resample(X_train_clv, y_train_clv)
```

Modeling

In [18]:

```
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline
from sklearn.linear_model import SGDClassifier

'''Build pipeline of classifiers'''
random_state = 11
# set all CPU
n_jobs = -1
# LogisticRegression
pipe_lr = Pipeline([('lr', LogisticRegression(random_state=random_state, n_jobs=n_jobs, ma
# RandomForestClassifier
pipe_rf = Pipeline([('rf', RandomForestClassifier(random_state=random_state, oob_score=True
# KNeighborsClassifier
pipe_knn = Pipeline([('knn', KNeighborsClassifier(n_jobs=n_jobs))])
# DecisionTreeClassifier
pipe_dt = Pipeline([('dt', DecisionTreeClassifier(random_state=random_state, max_features=
# BaggingClassifier
# note we use SGDClassifier as classier inside BaggingClassifier
pipe_bag = Pipeline([('bag', BaggingClassifier(base_estimator=SGDClassifier(random_state=ra
random_state=random_state, oob_score=True, n_
```

In [19]:

```
'''Set parameters for Grid Search '''
# set number
cv = StratifiedKFold(shuffle=True, n_splits=5, random_state=random_state)
# set for LogisticRegression
grid_params_lr = [{
    'lr_penalty': ['l2'],
    'lr_C': [0.3, 0.6, 0.7],
    'lr_solver': ['sag']
}]
# set for RandomForestClassifier
grid_params_rf = [{
    'rf_criterion': ['entropy'],
    'rf_min_samples_leaf': [80, 100],
    'rf_max_depth': [25, 27],
    'rf_min_samples_split': [3, 5],
    'rf_n_estimators': [60, 70]
}]
# set for KNeighborsClassifier
grid_params_knn = [{'knn_n_neighbors': [16, 17, 18]}]
# set for DecisionTreeClassifier
grid_params_dt = [{
    'dt_max_depth': [8, 10],
    'dt_min_samples_leaf': [1, 3, 5, 7]
}]
# set for BaggingClassifier
grid_params_bag = [{'bag_n_estimators': [20, 25, 30]}]
```

In [20]:

```
'''Grid search objects'''
# for LogisticRegression
gs_lr = GridSearchCV(pipe_lr, param_grid=grid_params_lr,
                     scoring='accuracy', cv=cv)
# for RandomForestClassifier
gs_rf = GridSearchCV(pipe_rf, param_grid=grid_params_rf,
                     scoring='accuracy', cv=cv)
# for KNeighborsClassifier
```



```

gs_knn = GridSearchCV(pipe_knn, param_grid=grid_params_knn,
                      scoring='accuracy', cv=cv)
# for DecisionTreeClassifier
gs_dt = GridSearchCV(pipe_dt, param_grid=grid_params_dt,
                    scoring='accuracy', cv=cv)
# for BaggingClassifier
gs_bag = GridSearchCV(pipe_bag, param_grid=grid_params_bag,
                    scoring='accuracy', cv=cv)

```

In [21]:

```

# models that we iterate over
look_for = [gs_lr, gs_rf, gs_knn, gs_dt, gs_bag]
# dict for later use
model_dict = {0:'Logistic_reg', 1:'RandomForest', 2:'Knn', 3:'DesionTree', 4:'Bagging with

```

In [22]:

```

''' Function to iterate over models and obtain results'''
# set empty dicts and list
result_acc = {}
result_auc = {}
models = []

for index, model in enumerate(look_for):
    print('+++++++ Start New Model ++++++')
    print('Estimator is {}'.format(model_dict[index]))
    model.fit(X_SMOTE_clv, y_SMOTE_clv)
    print('-----')
    print('best params {}'.format(model.best_params_))
    print('best score is {}'.format(model.best_score_))
    auc = roc_auc_score(y_test_clv, model.predict_proba(X_test_clv)[:,-1])
    print('-----')
    print('Test Set')
    print('-----')
    print('ROC_AUC is {} and accuracy rate is {}'.format(auc, model.score(X_test_clv,
    print('-----')
    print('Train Set')
    auc_train = roc_auc_score(y_SMOTE_clv, model.predict_proba(X_SMOTE_clv)[:,-1])
    print('ROC_AUC_train is {} and accuracy rate is {}'.format(auc_train, model.score
    print('-----')

    print('+++++++ End Model ++++++')
    models.append(model.best_estimator_)
    result_acc[index] = model.best_score_
    result_auc[index] = auc

```

```

+++++++ Start New Model ++++++
Estimator is Logistic_reg
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\linear_model\_sag.py:352: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
-----
best params {'lr_C': 0.3, 'lr_penalty': 'l2', 'lr_solver': 'sag'}
best score is 0.6708014667994657
-----
Test Set
-----
ROC_AUC is 0.7180303743819166 and accuracy rate is 0.6514161220043573
-----
Train Set
ROC_AUC_train is 0.7326940891370924 and accuracy rate is 0.6723033820292176
-----
+++++++ End Model ++++++

```


[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted w
ith feature names
    warnings.warn(
-----
best params {'knn__n_neighbors': 16}
best score is 0.792475416887623
-----
Test Set
-----
ROC_AUC is 0.6550829997645397 and accuracy rate is 0.6833696441539578
-----
Train Set
ROC_AUC_train is 0.924487752487394 and accuracy rate is 0.8265959575745447
-----
+++++++ End Model ++++++
+++++++ Start New Model ++++++
Estimator is DesionTree
-----
best params {'dt__max_depth': 10, 'dt__min_samples_leaf': 3}
best score is 0.7474487994748126
-----
Test Set
-----
ROC_AUC is 0.6029932893807393 and accuracy rate is 0.654320987654321
-----
Train Set
ROC_AUC_train is 0.8741798010565395 and accuracy rate is 0.7889733840304183
-----
+++++++ End Model ++++++
+++++++ Start New Model ++++++
Estimator is Bagging with SGDClassifier
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:706: UserWarning: Some inputs do not have OOB scores. This probably means too fe
w estimators were used to compute any reliable oob estimates.
    warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:712: RuntimeWarning: invalid value encountered in true_divide
    oob_decision_function = predictions / predictions.sum(axis=1)[:, np.newaxis]
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:706: UserWarning: Some inputs do not have OOB scores. This probably means too fe
w estimators were used to compute any reliable oob estimates.
    warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:712: RuntimeWarning: invalid value encountered in true divide
```



```

    oob_decision_function = predictions / predictions.sum(axis=1)[:, np.newaxis]
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:706: UserWarning: Some inputs do not have OOB scores. This probably means too fe
w estimators were used to compute any reliable oob estimates.
    warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:712: RuntimeWarning: invalid value encountered in true_divide
    oob_decision_function = predictions / predictions.sum(axis=1)[:, np.newaxis]
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:706: UserWarning: Some inputs do not have OOB scores. This probably means too fe
w estimators were used to compute any reliable oob estimates.
    warn(
C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\ensemble\_b
agging.py:712: RuntimeWarning: invalid value encountered in true_divide
    oob_decision_function = predictions / predictions.sum(axis=1)[:, np.newaxis]
-----
best params {'bag__n_estimators': 25}
best score is 0.6481902612968147
-----
Test Set
-----
ROC_AUC is 0.7312308688485989 and accuracy rate is 0.6855482933914306
-----
Train Set
ROC_AUC_train is 0.7215105929677479 and accuracy rate is 0.6534920952571542
-----
+++++++ End Model ++++++

```

In [23]:

```

predicted_y = []
expected_y = []
predictions = gs_bag.predict(X_test_clv.to_numpy())
predicted_y.extend(predictions)
expected_y.extend(y_test_clv)
report_test = classification_report(expected_y, predicted_y)
print('test set')
print(report_test)

```

```

C:\Users\S540\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:44
1: UserWarning: X does not have valid feature names, but BaggingClassifier was fitted with
feature names

```

```

warnings.warn(
test set

```

	precision	recall	f1-score	support
0	0.94	0.69	0.80	1240
1	0.19	0.64	0.29	137
accuracy			0.69	1377
macro avg	0.56	0.66	0.54	1377
weighted avg	0.87	0.69	0.75	1377

The best model is SGDClassifier The AUC train = 0.721 The AUC test = 0.731